

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report,
confusion_matrix

```

#Step 1: Data Preprocessing

```
data = pd.read_csv("Employee-Attrition.csv")
```

Explore the dataset

```
print(data.head())
```

```
print(data.info())
```

	Age	Attrition	BusinessTravel	DailyRate	Department
0	41	Yes	Travel_Rarely	1102	Sales
1	49	No	Travel_Frequently	279	Research & Development
2	37	Yes	Travel_Rarely	1373	Research & Development
3	33	No	Travel_Frequently	1392	Research & Development
4	27	No	Travel_Rarely	591	Research & Development

EmployeeNumber	DistanceFromHome	Education	EducationField	EmployeeCount
0	1	2	Life Sciences	1
1				
1	8	1	Life Sciences	1
2				
2	2	2	Other	1
4				
3	3	4	Life Sciences	1
5				
4	2	1	Medical	1
7				

	...	RelationshipSatisfaction	StandardHours	StockOptionLevel	\
0	...		1	80	0
1	...		4	80	1
2	...		2	80	0
3	...		3	80	0

4	...	4	80	1
	TotalWorkingYears	TrainingTimesLastYear	WorkLifeBalance	
YearsAtCompany \				
0	8	0	1	
6				
1	10	3	3	
10				
2	7	3	3	
0				
3	8	3	3	
8				
4	6	3	3	
2				
	YearsInCurrentRole	YearsSinceLastPromotion	YearsWithCurrManager	
0	4	0	5	
1	7	1	7	
2	0	0	0	
3	7	3	0	
4	2	2	2	

[5 rows x 35 columns]

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1470 entries, 0 to 1469

Data columns (total 35 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Age	1470 non-null	int64
1	Attrition	1470 non-null	object
2	BusinessTravel	1470 non-null	object
3	DailyRate	1470 non-null	int64
4	Department	1470 non-null	object
5	DistanceFromHome	1470 non-null	int64
6	Education	1470 non-null	int64
7	EducationField	1470 non-null	object
8	EmployeeCount	1470 non-null	int64
9	EmployeeNumber	1470 non-null	int64
10	EnvironmentSatisfaction	1470 non-null	int64
11	Gender	1470 non-null	object
12	HourlyRate	1470 non-null	int64
13	JobInvolvement	1470 non-null	int64
14	JobLevel	1470 non-null	int64
15	JobRole	1470 non-null	object
16	JobSatisfaction	1470 non-null	int64
17	MaritalStatus	1470 non-null	object
18	MonthlyIncome	1470 non-null	int64
19	MonthlyRate	1470 non-null	int64
20	NumCompaniesWorked	1470 non-null	int64
21	Over18	1470 non-null	object

22	OverTime	1470	non-null	object
23	PercentSalaryHike	1470	non-null	int64
24	PerformanceRating	1470	non-null	int64
25	RelationshipSatisfaction	1470	non-null	int64
26	StandardHours	1470	non-null	int64
27	StockOptionLevel	1470	non-null	int64
28	TotalWorkingYears	1470	non-null	int64
29	TrainingTimesLastYear	1470	non-null	int64
30	WorkLifeBalance	1470	non-null	int64
31	YearsAtCompany	1470	non-null	int64
32	YearsInCurrentRole	1470	non-null	int64
33	YearsSinceLastPromotion	1470	non-null	int64
34	YearsWithCurrManager	1470	non-null	int64

dtypes: int64(26), object(9)

memory usage: 402.1+ KB

None

Check for missing values

```
print(data.isnull().sum())
```

Explore the dataset

```
print(data.describe())
```

Data Visualization

```
sns.countplot(x='Attrition', data=data)
```

```
plt.show()
```

Age	0
Attrition	0
BusinessTravel	0
DailyRate	0
Department	0
DistanceFromHome	0
Education	0
EducationField	0
EmployeeCount	0
EmployeeNumber	0
EnvironmentSatisfaction	0
Gender	0
HourlyRate	0
JobInvolvement	0
JobLevel	0
JobRole	0
JobSatisfaction	0
MaritalStatus	0
MonthlyIncome	0
MonthlyRate	0
NumCompaniesWorked	0
Over18	0
OverTime	0

```

PercentSalaryHike      0
PerformanceRating      0
RelationshipSatisfaction  0
StandardHours          0
StockOptionLevel       0
TotalWorkingYears      0
TrainingTimesLastYear  0
WorkLifeBalance        0
YearsAtCompany          0
YearsInCurrentRole     0
YearsSinceLastPromotion 0
YearsWithCurrManager   0
dtype: int64

```

```

      Age      DailyRate  DistanceFromHome      Education
EmployeeCount \
count  1470.000000    1470.000000      1470.000000    1470.000000
1470.0
mean    36.923810     802.485714         9.192517     2.912925
1.0
std      9.135373     403.509100         8.106864     1.024165
0.0
min     18.000000     102.000000         1.000000     1.000000
1.0
25%     30.000000     465.000000         2.000000     2.000000
1.0
50%     36.000000     802.000000         7.000000     3.000000
1.0
75%     43.000000    1157.000000        14.000000     4.000000
1.0
max     60.000000    1499.000000        29.000000     5.000000
1.0

```

```

      EmployeeNumber  EnvironmentSatisfaction      HourlyRate
JobInvolvement \
count    1470.000000      1470.000000    1470.000000
1470.000000
mean    1024.865306         2.721769     65.891156
2.729932
std      602.024335         1.093082     20.329428
0.711561
min       1.000000         1.000000     30.000000
1.000000
25%      491.250000         2.000000     48.000000
2.000000
50%     1020.500000         3.000000     66.000000
3.000000
75%     1555.750000         4.000000     83.750000
3.000000
max     2068.000000         4.000000    100.000000

```

4.000000


	JobLevel	...	RelationshipSatisfaction	StandardHours	\
count	1470.000000	...	1470.000000	1470.0	
mean	2.063946	...	2.712245	80.0	
std	1.106940	...	1.081209	0.0	
min	1.000000	...	1.000000	80.0	
25%	1.000000	...	2.000000	80.0	
50%	2.000000	...	3.000000	80.0	
75%	3.000000	...	4.000000	80.0	
max	5.000000	...	4.000000	80.0	

	StockOptionLevel	TotalWorkingYears	TrainingTimesLastYear	\
count	1470.000000	1470.000000	1470.000000	
mean	0.793878	11.279592	2.799320	
std	0.852077	7.780782	1.289271	
min	0.000000	0.000000	0.000000	
25%	0.000000	6.000000	2.000000	
50%	1.000000	10.000000	3.000000	
75%	1.000000	15.000000	3.000000	
max	3.000000	40.000000	6.000000	

	WorkLifeBalance	YearsAtCompany	YearsInCurrentRole	\
count	1470.000000	1470.000000	1470.000000	
mean	2.761224	7.008163	4.229252	
std	0.706476	6.126525	3.623137	
min	1.000000	0.000000	0.000000	
25%	2.000000	3.000000	2.000000	
50%	3.000000	5.000000	3.000000	
75%	3.000000	9.000000	7.000000	
max	4.000000	40.000000	18.000000	

	YearsSinceLastPromotion	YearsWithCurrManager
count	1470.000000	1470.000000
mean	2.187755	4.123129
std	3.222430	3.568136
min	0.000000	0.000000
25%	0.000000	2.000000
50%	1.000000	3.000000
75%	3.000000	7.000000
max	15.000000	17.000000

[8 rows x 26 columns]



```
# Encode categorical variables
label_encoder = LabelEncoder()
categorical_columns = data.select_dtypes(include=['object']).columns

for column in categorical_columns:
    data[column] = label_encoder.fit_transform(data[column])

# Split data into features (X) and target (y)
X = data.drop('Attrition', axis=1)
y = data['Attrition']

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_state=42)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

#Step 2: Model Building
# Initialize and train the logistic regression model
logistic_regression_model = LogisticRegression()
logistic_regression_model.fit(X_train, y_train)

LogisticRegression()

# Initialize and train the decision tree classifier
decision_tree_model = DecisionTreeClassifier()
decision_tree_model.fit(X_train, y_train)
```

```

DecisionTreeClassifier()

# Initialize and train the random forest classifier
random_forest_model = RandomForestClassifier()
random_forest_model.fit(X_train, y_train)

RandomForestClassifier()

#Step 3: Calculate Performance Metrics
# Predict using logistic regression model
y_pred_lr = logistic_regression_model.predict(X_test)

# Calculate accuracy
accuracy_lr = accuracy_score(y_test, y_pred_lr)

# Generate classification report
classification_report_lr = classification_report(y_test, y_pred_lr)

# Generate confusion matrix
confusion_matrix_lr = confusion_matrix(y_test, y_pred_lr)

# Predict using decision tree model
y_pred_dt = decision_tree_model.predict(X_test)

# Calculate accuracy
accuracy_dt = accuracy_score(y_test, y_pred_dt)

# Generate classification report
classification_report_dt = classification_report(y_test, y_pred_dt)

# Generate confusion matrix
confusion_matrix_dt = confusion_matrix(y_test, y_pred_dt)

# Predict using random forest model
y_pred_rf = random_forest_model.predict(X_test)

# Calculate accuracy
accuracy_rf = accuracy_score(y_test, y_pred_rf)

# Generate classification report
classification_report_rf = classification_report(y_test, y_pred_rf)

# Generate confusion matrix
confusion_matrix_rf = confusion_matrix(y_test, y_pred_rf)

#Step 4: Display Results
# Display results for Logistic Regression
print("Logistic Regression Metrics:")
print("Accuracy:", accuracy_lr)
print("Classification Report:\n", classification_report_lr)
print("Confusion Matrix:\n", confusion_matrix_lr)

```

```
# Display results for Decision Tree
print("\nDecision Tree Metrics:")
print("Accuracy:", accuracy_dt)
print("Classification Report:\n", classification_report_dt)
print("Confusion Matrix:\n", confusion_matrix_dt)

# Display results for Random Forest
print("\nRandom Forest Metrics:")
print("Accuracy:", accuracy_rf)
print("Classification Report:\n", classification_report_rf)
print("Confusion Matrix:\n", confusion_matrix_rf)
```

Logistic Regression Metrics:

Accuracy: 0.891156462585034

Classification Report:

	precision	recall	f1-score	support
0	0.91	0.98	0.94	255
1	0.68	0.33	0.45	39
accuracy			0.89	294
macro avg	0.79	0.65	0.69	294
weighted avg	0.88	0.89	0.87	294

Confusion Matrix:

```
[[249  6]
 [ 26 13]]
```

Decision Tree Metrics:

Accuracy: 0.7857142857142857

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.88	0.88	255
1	0.18	0.18	0.18	39
accuracy			0.79	294
macro avg	0.53	0.53	0.53	294
weighted avg	0.78	0.79	0.78	294

Confusion Matrix:

```
[[224 31]
 [ 32  7]]
```

Random Forest Metrics:

Accuracy: 0.8775510204081632

Classification Report:

	precision	recall	f1-score	support
0	0.88	1.00	0.93	255

1	0.80	0.10	0.18	39
accuracy			0.88	294
macro avg	0.84	0.55	0.56	294
weighted avg	0.87	0.88	0.83	294
Confusion Matrix:				
[[254 1]				
[35 4]]				